

Computer Arithmetic

Signed and Unsigned

Topics

Overflow

Addition and carry, subtraction and borrow

Processor flags: Z, N, C, V

2's complement representation of negative numbers


Comparisons


Implicit type conversions (C craziness)

9 9 9 9 9 9

9 9 9 9 9 9

0 0 0 0 0 0





What Happens at Overflow?

Grade School Addition

247

9

Grade School Addition

$$\begin{array}{r} 1 \\ 247 \\ 9 \\ \hline 6 \end{array}$$

Grade School Addition

$$\begin{array}{r} 1 \\ 247 \\ 9 \\ \hline 56 \end{array}$$

Grade School Addition

$$\begin{array}{r} 1 \\ 247 \\ 09 \\ - - - \\ 256 \end{array}$$

Addition (Hexadecimal)

F7

09

- - -

Addition (Hexadecimal)

1

F7

09

0

Addition (Hexadecimal)

11

F7

09

00

Addition (Hexadecimal)

11

F7

09

100

Addition (Binary)

11110111

00001001

Addition (Binary)

$$\begin{array}{r} 1 \\ 11110111 \\ 00001001 \\ \hline 0 \end{array}$$

Addition (Binary)

```
      11
  11110111
  00001001
  -----
      00
```

Addition (Binary)

```
      111
    11110111
    00001001
    -----
           000
```

Addition (Binary)

```
      1111
    11110111
    00001001
    -----
      0000
```

Addition (Binary)

```
    11111
  11110111
  00001001
  -----
    00000
```

Addition (Binary)

```
  111111
 11110111
00001001
-----
 000000
```

Addition (Binary)

1111111

11110111

00001001

0000000

Addition (Binary)

11111111

11110111

00001001

00000000

Addition (Binary)

```
11111111
 11110111
 00001001
-----
00000000
```

Result: 00000000 (only room for 8-bits)
Carry (C): 1 (extra bit)

To hold the result of adding two n-bit numbers requires n+1 bits

add32/

add32

00000000 + 00000000 = 00000000 : Z=1, C=0

00000000 + 00000001 = 00000001 : Z=0, C=0

00000001 + 00000001 = 00000002 : Z=0, C=0

00000002 + 00000001 = 00000003 : Z=0, C=0

...


fffffffffe + 00000001 = ffffffff : Z=0, C=0

ffffffff + 00000001 = 00000000 : Z=1, C=1

...

ffffffff + ffffffff = ffffffffe : Z=0, C=1

CPSR



Arithmetic instructions set Z, C (N, V)

Logic instructions just set Z (N)

(We will cover N soon, and V later)

```
// Multiple precision addition  
// http://godbolt.org/g/HMYrme
```

```
int64_t add64(int64_t a, int64_t b)  
{  
    return a + b;  
}
```

```
add64:  
    adds r0, r0, r2  
    adcs r1, r1, r3  
    bx    lr
```

Negative (signed) Numbers

Up to now, all binary numbers have been positive (unsigned)


How would you define negative (signed) numbers?

A clever way of defining -1 is to say that -1 is the number that when added to 1, results in 0 (mod 16)

$$0xf + 0x1 = 0x10 \% 16 = 0x0$$

0xf can be *interpreted* as -1

Note that modulo addition of signed numbers is exactly the same as addition of unsigned numbers!!



For 4-bit numbers,

0x0 = 0

0xf = -1

0xe = -2

...

0x8 = -8 (could be interpreted as 8)

0x7 = 7

...

0x1 = 1

0x0 = 0

if we choose to *interpret* 0x8 as -8,
then the most-significant bit of the
number indicates that it is negative (n)

What is the maximum value for a int32?



606

integer



share improve this question



92

I can never remember that number. I need a memory rule.

edited May 28 '14 at 14:09



Ben Hoffstein

49.5k ● 5 ● 66 ● 101

asked Sep 18 '08 at 17:18



Flinkman

5,181 ● 4 ● 18 ● 48

107 Why would you need the exact number? I remember " $(2^{31})-1$ " or " ± 2 billion" and that's good enough for everything I ever needed. — Joachim Sauer Mar 3 '09 at 11:21

27 unsigned: $2^{32}-1 = 4 \cdot 1024^3 - 1$; signed: $-2^{31} \dots +2^{31}-1$, because the sign-bit is the highest bit. Just learn $2^0=1$ to $2^{10}=1024$ and combine. $1024=1k$, $1024^2=1M$, $1024^3=1G$ — comonad Mar 28 '11 at 20:01

6 I generally remember that every 3 bits is about a decimal digit. This gets me to the right order of magnitude: 32 bits is 10 digits. — Barnar Oct 2 '13 at 15:11

30 Answers

active

oldest

votes



2397

share improve this answer



It's 2,147,483,647. Easiest way to memorize it is via a tattoo.

edited Oct 20 '14 at 16:30



Allbite

1,415 ● 1 ● 13 ● 15

answered Sep 18 '08 at 17:20



Ben Hoffstein

49.5k ● 5 ● 66 ● 101

How do we negate an 8-bit number?

Subtract the number from 100000000!

```
11111111
10000000
 00000001
-----
11111111
```

11111111 (0xff) is -1

The 1's on the top are borrows

Another way

Rewrite 100000000 as 11111111 + 1

$$-x = (11111111+1)-x = (11111111-x)+1 = \sim x+1$$

11111111

-00000001

$$11111110 = \sim 00000001 \text{ (~ means flip bits)}$$

11111110

+00000001

11111111

Subtraction is converted to

$$a - b = a + \sim b + 1$$

$$01 - 00 = 01 + ff + 01 = 01 + c$$

$$01 - 01 = 01 + fe + 01 = 00 + c$$

$$01 - 02 = 01 + fd + 01 = ff$$

Why?

- Only one arithmetic op: add
- Need to think like this
to understand when a carry is generated

sub32/

00000000-00000000 = 00000000+00000000 = 00000000

00000000-00000001 = 00000000+ffffffff = ffffffff

ffffffff-00000001 = ffffffff+ffffffff = fffffffe

fffffffe-00000001 = fffffffe+ffffffff = fffffffd

...

00000000-ffffffff = 00000000+00000001 = 00000001

00000001-ffffffff = 00000001+00000001 = 00000002

00000001-fffffffd = 00000001+00000003 = 00000004

...

ffffffff-ffffffff = ffffffff+00000001 = 00000000

ffffffff-fffffffd = ffffffff+00000003 = 00000002

// What happens if the timer overflows?

```
unsigned timer_get_time(void) {  
    return GET32(SYSTIMERCLO);  
}
```

```
void delay_us(unsigned us) {  
    unsigned rb = timer_get_time();  
    while (1) {  
        unsigned ra = timer_get_time();  
        // subtraction results in a positive  
        // value even if the timer overflows  
        if ((ra - rb) >= us) {  
            break;  
        }  
    }  
}
```

**Addition / subtraction of
signed and unsigned numbers
is the same!**

Comparison

unsigned comparison: ucmp32/

>

00000001-00000000=00000001: C=1

ffffffff-00000001=ffffffffff: C=1

ffffffff-00000000=ffffffff: C=1

==

00000000-00000000=00000000: C=1

00000001-00000001=00000000: C=1

ffffffff-ffffffff=00000000: C=1


<

00000000-00000001=ffffffff: C=0

00000000-ffffffff=00000001: C=0



00000001-ffffffff=00000002: C=0


uge = c




- In two's complement, when you exceed the maximum value of int (2,147,483,647), you “wrap around” to negative numbers:



PSY - GANGNAM STYLE (강남스타일) M/V

 officialpsy 

 7,600,830



-2142584554


+ Add to  Share *** More

 8,761,309  1,139,933


- Here is the link after Google upgraded to 64-bit integers:



PSY - GANGNAM STYLE (강남스타일) M/V


 officialpsy 

 7,600,830

2,152,382,740

+ Add to  Share *** More

 8,761,309  1,139,933



signed comparison cmp32/

>

7fffffff - 00000000 = 7fffffff: N=0, V=0

00000000 - 80000000 = 80000000: N=1, V=1

7fffffff - 80000000 = ffffffff: N=1, V=1

==

00000000 - 00000000 = 00000000: N=0, V=0

80000000 - 80000000 = 00000000: N=0, V=0

7fffffff - 7fffffff = 00000000: N=0, V=0

<

00000000 - 7fffffff = 80000001: N=1, V=0

80000000 - 00000000 = 80000000: N=1, V=0

80000000 - 7fffffff = 00000001: N=0, V=1

overflow occurs if the sign of the two numbers being compared are the same, and the sign of the result is different $(n1==n2) \ \& \ (n!=n1)$

signed comparison cmp32/

>

7fffffff - 00000000 = 7fffffff: N=0, V=0

00000000 - 80000000 = 80000000: N=1, V=1

7fffffff - 80000000 = ffffffff: N=1, V=1

==

00000000 - 00000000 = 00000000: N=0, V=0

80000000 - 80000000 = 00000000: N=0, V=0

7fffffff - 7fffffff = 00000000: N=0, V=0

<

00000000 - 7fffffff = 80000001: N=1, V=0

80000000 - 00000000 = 80000000: N=1, V=0

80000000 - 7fffffff = 00000001: N=0, V=1

bge: signed greater than or equal (n == v)

blt: signed less than (n != v)

```
int ge() { return !v ? !n : n }
```

Code	Suffix	Flags	Meaning
0000	EQ	Z set	equal
0001	NE	Z clear	not equal
0010	CS	C set	unsigned higher or same
0011	CC	C clear	unsigned lower
0100	MI	N set	negative
0101	PL	N clear	positive or zero
0110	VS	V set	overflow
0111	VC	V clear	no overflow
1000	HI	C set and Z clear	unsigned higher
1001	LS	C clear or Z set	unsigned lower or same
1010	GE	N equals V	greater or equal
1011	LT	N not equal to V	less than
1100	GT	Z clear AND (N equals V)	greater than
1101	LE	Z set OR (N not equal to V)	less than or equal
1110	AL	(ignored)	always

**Comparison of
signed and unsigned numbers
is NOT the same!**

Type Conversion

Jedi Job

Interview Questions

```
uint16_t x = 0xffff;
```

```
uint32_t y = x;
```

```
// what is the value of y?
```

Type Hierarchy

uint32



uint16



uint8

Types are *sets* of allowed values

Arrow indicate *subsets*: $\text{uint8} \subset \text{uint16}$

uint32



uint16



uint8

Safe (values preserved)

```
int16_t x = -1;
```

```
int32_t y = x;
```

```
// what is the value of y?
```

int32



int16



int8

Safe (values preserved)

```
// Sign extend
```

int8 0xFE -> int32 0xFFFFFFFF

int8 0x7E -> int32 0x0000007E

// Assembly language

LSL $r_0, r_0, \#24$

ASR r0,r0,#24

[illegible]

asr

The diagram shows a horizontal row of 24 square boxes representing bits. The first 8 boxes from the left contain the value '1', and the remaining 16 boxes contain the value '0'. To the left of the first box is the label 'asr' in bold black font. Above the first two boxes is a large black curved arrow pointing downwards and to the right, indicating a right shift operation.

[illegible]


```
int32_t x = -1;
```

```
int16_t y = x;
```

```
// what is the value of y?
```

int32



int16



int8

Defined (throw away msb)

Dangerous (doesn't preserve all values)

```
int32_t  x = -1;  
uint32_t y =  x;
```

```
// what is the value of y?
```

uint32 ← int32

uint16 ← int16

uint8 ← int8

**Defined
(copies bits*)**

uint32 ← int32

uint16 ← int16

uint8 ← int8

Dangerous! (neg become large)

```
uint32_t x = 0xffffffff;
```

```
int32_t y = x;
```

```
// what is the value of y?
```

uint32 → int32

uint16 → int16


uint8 → int8

**Technically Not Defined
(copies bits)**

```
uint16_t x = 0xffff;
```

```
int32_t y = x;
```

```
// what is the value of y?
```

Safe (values preserved)

```
int16_t x = -1;  
uint32_t y = x;
```

```
// what is the value of y?
```

uint32

int32



uint16

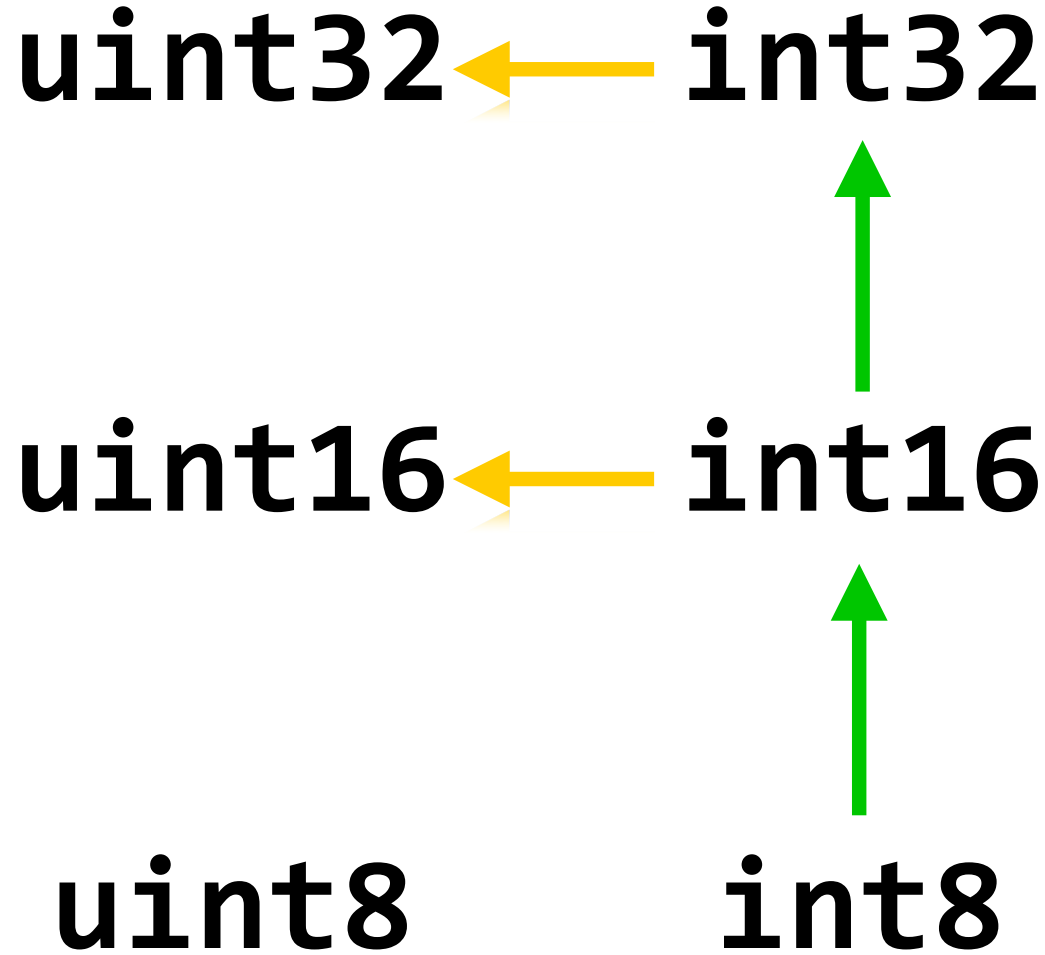
int16

uint8

int8

??





Defined, Dangerous

The semantics of numeric casts are:

Casting from a larger integer to a smaller integer (e.g. u32 -> u8) will truncate

Casting from a smaller integer to a larger integer (e.g. u8 -> u32) will zero-extend if the source is unsigned sign-extend if the source is signed


Casting between two integers of the same size (e.g. i32 -> u32) is a no-op

Binary Operations


Implicit Type Conversion (Coercion)

Resulting type from a binary operator (ARM32)

	u8	u16	u32	u64	i8	i16	i32	i64
u8	i32	i32	u32	u64	i32	i32	i32	i64
u16	i32	i32	u32	u64	i32	i32	i32	i64
u32	u32	u32	u32	u64	u32	u32	u32	i64
u64	u64	u64	u64	u64	u64	u64	u64	u64
i8	i32	i32	u32	u64	i32	i32	i32	i64
i16	i32	i32	u32	u64	i32	i32	i32	i64
i32	i32	i32	u32	u64	i32	i32	i32	i64
i64	i64	i64	i64	u64	i64	i64	i64	i64



Bugs, Bugs, Bugs




```
#include <stdio.h>
```

```
int main(void)
```

```
{
```

```
    int a = -20;
```

```
    unsigned int b = 6;
```

```
    if( a < b )
```

```
        printf("-20 < 6 - all is well  
\n");
```

```
    else
```

```
        printf("-20 >= 6 - omg \n");
```

```
}
```

**Whenever you mix
signed and unsigned numbers
you get in trouble**

Bjarne Stroustrup

C Promotion Rules

6.3.1.3 Signed and unsigned integers conversions

1 When a value with integer type is converted to another integer type, if the value can be represented by the new type, it is unchanged.

2 Otherwise, if the new type is unsigned, the value is converted by repeatedly adding or subtracting one more than the maximum value that can be represented in the new type until the value is in the range of the new type.

3 Otherwise, if the new type is signed and the value cannot be represented in it; either the result is implementation-defined or an implementation-defined signal is raised.

6.3.1.8 Usual arithmetic conversions

1 If both operands have the same type, then no further conversion is needed.

2 Otherwise, if both operands have signed integer types or both have unsigned integer types, the operand with the type of lesser integer conversion rank is converted to the type of the operand with greater rank.

3 Otherwise, if the operand that has unsigned integer type has rank greater or equal to the rank of the type of the other operand, then the operand with signed integer type is converted to the type of the operand with unsigned integer type.

4 Otherwise, if the type of the operand with signed integer type can represent all of the values of the type of the operand with unsigned integer type, then the operand with unsigned integer type is converted to the type of the operand with signed integer type.

5 Otherwise, both operands are converted to the unsigned integer type corresponding to the type of the operand with signed integer type.